

Tomorrow's Newspaper

The making of EON, a machine that reads annual reports, and an honest test of what it could see

Gabriel George

Technical white paper

30 September 2026 · Edition 4

5 versions, from a folder of scripts to a system
29,375 answers from a language model, each one dated
1 question: could it see ahead?

I built a machine to read every annual report in the market. This is how it was built,

what it read, and whether it had simply read tomorrow's newspaper.

Contents

Abstract	1
1. A reader for every annual report	1
Part I · Building the reader	3
2. The workshop: learning from great companies	3
3. From scripts to a processor	4
4. Fintel: an application	4
5. EON: one question, every filing	5
6. Workflows as hypotheses	12
Part II · What it read	13
7. Every question it was asked	13
8. Was it right?	14
Part III · What it taught	22
9. How the model went wrong	22
10. Why it was worth doing	23
11. The bet	24
12. Conclusion	25
Appendix A: methods	26
Appendix B: related work and references	26
Appendix C: reproduction	27

Abstract

Over seventeen months I built five versions of the same idea: a machine that reads company annual reports the way a patient analyst would, and writes down what it thinks. It began in May 2025 as a folder of scripts studying thirty years of great companies. It became EON: a system of downloaders, headless browsers, key rotation, file locks, leases, heartbeats, a SQLite database, a command line, a web interface and a Discord alarm, all built around one constraint, a daily quota of model requests. Together the versions stored 29,375 answers from a language model, including 6,973 three-lens readings of 10-K filings, each ending in a single verdict.

This paper is mostly about how that machine was built and why it needed every part. It is also about what the machine read, and whether it was right. Graded naively, its BUY calls beat its SELL calls by 13 percentage points a year. But the model had been trained on the years it was being graded on: it may have read tomorrow's newspaper. Graded only on what it could not have known, the edge shrinks; its SELL calls stop working; a score written down in advance pointed the wrong way; and an options scan turned out to see how far a stock would move, but not which way.

It was never going to rival Citadel, Point72 or the Medallion Fund. It was worth building anyway, as a finance project, an AI project and a systems project at once. The paper ends with a sealed bet: 1,257 verdicts frozen at publication, to be graded in October 2027.

Research, not investment advice.

1. A reader for every annual report

Every public company in the United States files a 10-K each year. It is the fullest account a company gives of itself: what it sells, how it makes money, what could go wrong, what it owes. The median 10-K in this project runs to 141 pages. Thousands are filed every spring. Professional analysts read the few they are paid to cover; almost nobody reads the rest.


I wanted a reader that never tires. Give it any filing and it would read the whole thing, ask the same careful questions every time, and write down what it concluded, so that I could come back a year later and see whether it had been right. A language model can now read 141 pages in a minute. The hard part is everything around that minute.

Figure 1 is what the finished machine produces. Its verdict on Apple reads: "SELL. Conviction: High. Apple is an undeniably excellent company, but a high conviction SELL rating is warranted given the significant disconnect between its robust fundamentals and its highly stretched valuation." It is specific, fluent and confident. Whether it is right is a separate question, and Part II is about how hard that question turned out to be.

The paper has three parts. Part I tells how the machine was built, version by version, and why each component exists. Part II describes everything it was asked and tests its answers against what happened. Part III is about what the project taught: where the model failed, and why the whole thing was worth doing.

One filing becomes one countable verdict

Specimen: Apple 10-K for fiscal 2025, filed 31 Oct 2025 / illustrative, not evidence



Six months later, Apple had trailed SPY by 0.9 points: a draw.

Source: data/eon.db, reading 1; evaluation/trades.csv. The reading's claims are not audited.


Figure 1: One filing becomes one countable verdict. EON's first stored reading: Apple's 10-K for fiscal 2025, filed 31 October 2025. The model reads the filing as Warren Buffett, Nassim Taleb and a contrarian would, then gives one verdict. Six months later Apple had trailed the S&P 500 by 0.9 points: a draw. Sources: data/eon.db; evaluation/trades.csv.

Part I · Building the reader

EON was built five times. Each version answered a failure of the last and kept what worked (Figure 2).

Five versions, one idea carried forward

May 2025 to June 2026 / each version kept what the last one taught



Sources: origin files and dates; git history of EON; data/archive/fintel.db; data/eon.db.

Figure 2: Five versions, one idea carried forward. From scripts studying great companies to an application that reads every filing. Dates come from the origin project's files and the repository's git history. Sources: the origin project; git history; data/archive/fintel.db; data/eon.db.

2. The workshop: learning from great companies

The project did not begin as a stock picker. It began, in May 2025, as a question about excellence: what do great companies have in common, and which other companies have it too?

The first answer was a folder of Python scripts called `10K_automator`. They downloaded up to thirty years of annual reports for 39 long-run compounders (Apple, Costco, Copart, Fair Isaac, O'Reilly, AutoZone and others), converted each one to PDF with a headless Chrome browser, extracted the text, and asked Gemini 2.5 Flash, then an April preview, to describe what made the company work that year. A second pass distilled each company's thirty readings into success factors; a third distilled all 39 into a meta-analysis.

It found six “universal” factors: operational excellence, capital allocation, continuous innovation, domain expertise, culture, and customer focus. The model credited continuous innovation to 100% of the companies. That should have been a warning. A factor that every great company has, and that nearly every annual report claims, cannot tell great companies from anyone else.

The scripts then read the recent filings of about 1,900 other companies and scored each from 0 to 100 on its resemblance to the pattern. Between 3 and 13 May they produced 11,741 readings of individual filings, 4,641 of them on 13 May alone. Two more experiments followed: a contrarian scan on 26 May, scoring 1,933 companies on six dimensions of hidden strength, and an options scan on 4 June that proposed 951 trades in calls or puts, run on ten API keys with ten progress files so that a crash would not lose the day.

Almost everything in EON started in those five weeks: reading whole filings through a browser, parallel workers, progress files that survive a crash, and the ambition to read the entire market. So did two problems. The outputs had no fixed format; a script written later to flatten them found the six success factors spelled about 130 different ways. And the model did not know what day it was: of 1,919 comparative analyses run in May 2025, 1,916 dated themselves 2024.

The outputs were saved and never changed. More than a year later, that turned out to be their most valuable property (Section 8.4).

3. From scripts to a processor

In October 2025 the scripts became a single reusable module, `standardized_sec_ai`. Its centre was `tenk_processor.py`, which downloaded, converted and read filings for any ticker, and its most important change was a single word: Pydantic.

Pydantic is a Python library for declaring the exact shape of data. Handed to Gemini as a response schema, it forces the model to answer in named fields of fixed types, and rejects any answer that does not fit. The project's own notes from that month celebrate it: “No more JSON parsing errors!” The 130 spellings of six factors could not happen again.

The same folder held `ppee.py`, the three-lens prompt from which EON's is still drawn: read the filing as Buffett would (moat, management, price), as Taleb would (fragility, tail risk, optionality), and as a contrarian would (what everyone believes, and why they might be wrong).

4. Fintel: an application

In December 2025 the processor grew into an application called Fintel, with a Streamlit web interface, a SQLite database and a visual workflow builder. A user could chain steps (pick companies, fetch filings, run a lens, filter, aggregate, export) and save the chain. The archive holds 16 of them: comparing three chip makers, reading three nuclear start-ups, pulling executive pay from a proxy statement.

The builder was good for discovering which questions mattered and bad for measuring anything. No two chains asked the same thing, so nothing accumulated. It was removed after three weeks.

What stayed was the infrastructure the builder had needed. Fintel was the first version that ran several analyses at once, and running several analyses at once broke things. Its first rate limiter used Python's `threading.Lock`, which only works inside one process. When the command line and the web interface ran together, or a batch spawned several processes, each had its own lock and the calls collided: the API answered with errors 429 and 503. The fix, recorded in a design note from January 2026, replaced it with a file lock (`portalocker`) that every process on the machine could see. The same month brought a nightly queue, an SEC rate limiter, and leases with heartbeats so that a crashed worker could not strand a company forever.


5. EON: one question, every filing

In February 2026 the project was renamed EON, the Erebus Observatory Network, and made one decisive change. Instead of many questions asked of a few companies, it would ask one question of every company: the three-lens prompt, with a 35-field schema, applied to every listed company's recent 10-Ks. Essays became a panel, the same measurement taken across 1,358 companies and six fiscal years. Everything in Part II depends on that.

Asking one question of every filing is easy to say. Figure 3 shows what it took.

Five layers between a filing and a verdict

EON's architecture / module names from the repository



Source: eon/ package structure; limits from eon/ai/api_config.py. Arrows: the path of one filing.

Figure 3: Five layers between a filing and a verdict. Each box is a module in the repository. Arrows trace one filing: fetched from EDGAR under the SEC's rate policy, printed to PDF by headless Chrome, extracted to text, sent with the schema through the key manager and request queue to Gemini, validated, and stored. Source: the eon/ package.

Each layer exists because something failed without it. The rest of this section takes them in turn.


5.1 Getting the words: why a PDF, and not XML?

The SEC publishes every filing on EDGAR, free, as HTML. EON fetches them with `sec-edgar-downloader` through its own queue, which keeps well inside the SEC's fair-access limit of ten requests a second: two seconds between requests, at most five at once, with a file lock so that parallel workers on one machine share the budget. Each filing is cached by ticker and year and never fetched twice; the cache now holds 27.5 GB of PDFs.

Why not simply read the HTML, or use the SEC's machine-readable XBRL data? The measurements answer both (Figure 4).

A 141-page filing becomes one word

The median 10-K on its way through EON / characters, logarithmic scale



Sources: evaluation/pipeline-measurements.json (17 raw 10-Ks; 40 PDFs); data/eon.db.

Figure 4: A 141-page filing becomes one word. Characters at each stage for the median 10-K: raw EDGAR HTML across 17 sampled filings, their visible text, the text EON extracts from its PDF (40 sampled filings), the stored reading, and the verdict. Sources: evaluation/pipeline-measurements.json; data/eon.db.

Raw EDGAR HTML is 5.6 to 19.6 times larger than the text a person would read. Most of the difference is markup, and much of that is inline XBRL: between 1,170 and 3,888 machine-readable tags per 10-K, plus a hidden header of tagged data. In Kraft Heinz's 2022 10-K, the invisible XBRL header (645,000 characters) is larger than the entire visible report (642,000). Feed that HTML to a model and most of what you pay for is markup; strip the tags naively and the hidden data pours into the text.

XBRL itself is excellent for numbers: revenue, debt, cash flow, tagged and comparable. But EON's questions are about judgement: moats, fragility, what management says about risk. That lives in the narrative, which XBRL does not structure.

So EON prints each filing the way a person would see it. A headless Chrome browser opens the HTML, waits between 3 and 15 seconds depending on the file's size so that the page finishes rendering, and calls Chrome's own print-to-PDF command. The result is exactly what a reader sees: tables laid out, hidden XBRL gone, 141 pages at the median. PyPDF2 then extracts the text in about 6 seconds. Nothing is summarised or cut; a filing too long for the model's context is skipped, not truncated.

The route costs a browser per filing, and browsers misbehave. A conversion that hangs is abandoned after 180 seconds and the browser restarted. Orphaned Chrome processes accumulate over a long batch, so EON sweeps them up every 50 companies, or sooner if memory use passes 80%. Tables and layout can still be mangled in extraction; that path has not been audited, and parsing HTML directly would be cleaner. For a machine built to read like a person, printing the page like a person was the simplest reliable answer.

5.2 Asking the question

The median filing arrives as about 394,000 characters, roughly 100,000 tokens. EON appends all of it to the three-lens prompt and sends one request to Gemini 2.5 Flash, with the 35-field schema attached. The model returns about 30,000 characters, 5,000 words of analysis, and one verdict (Figure 4).

Pydantic validates every response before it is stored. A response that does not fit is not a reading. Failures are sorted by cause before they are retried: a rate limit waits for exactly as long as the API's own retry-after field says; a transient server error backs off and tries again; a context overflow is marked skipped and never retried. Across all 6,973 three-lens readings, the model's input came to roughly 0.7 billion tokens.

5.3 Twenty requests a day

The binding constraint was not compute, storage or money. It was requests. The February configuration rotated 25 API keys, each allowed 20 requests a day, with a mandatory 65-second pause between requests on any one key. That is 500 filings a day, at most.

Every other design decision in EON follows from that number, and getting it right took three attempts.

The first, in Fintel, was the in-process lock that different processes could not see. The second was the file lock that fixed it: one lock for the whole machine, held through each request and its 65-second pause. It worked, and it made 25 keys exactly as slow as one. The third, in EON, gave every key its own lock file and put a 25-slot semaphore across all of them, so that 25 requests could run at once while no single key was ever used twice at the same moment. The pause adapts: it shrinks after a run of successes, down to about 21 seconds, and grows by half after a rate-limit error.

A fourth bug was quieter. On 6 February 2026 a commit titled "Remove double-counting of API key usage that halved effective capacity" found that every request was recorded twice, once by the model provider and again by seven of its callers, so each key reached its limit of 20 after 10 real requests. For a while the machine had been running at half speed and reporting itself full.

Parallelism is what turned 25 keys into 500 readings a day. Figure 5 shows the difference between the two ways the command line runs.


One company runs in a line; a batch runs in lanes

How the CLI executes / schematic, not to scale

A eon analyze AAPL --years 3



B eon batch companies.csv (25 keys; five lanes shown)




Source: eon/cli, eon/ui/services/batch_queue.py, eon/ai/request_queue.py.

Figure 5: One company runs in a line; a batch runs in lanes. A single analysis downloads, prints, extracts, asks and saves one fiscal year after another. A batch starts up to 25 workers with staggered starts; each leases one company, reserves one key, and renews its lease with a heartbeat. When every key is spent, all workers wait for the reset at midnight Pacific. Schematic. Sources: eon/cli; eon/ui/services/batch_queue.py; eon/ai/request_queue.py.

The market-wide batch ran from 8 to 21 February 2026 and read 6,568 filings from 1,327 companies (Figure 6). At its busiest it finished 189 readings in an hour, one every 19 seconds across all workers. But it was active for only 62 of the thirteen days' hours. For about 80% of its life, the fastest part of the machine was waiting for midnight.

The quota is the clock

Market-wide batch, 8 to 21 February 2026 / readings per hour, UTC



Nominal allowance: 500 requests/day.

Dotted lines: quota reset, midnight Pacific.

Source: data/eon.db, batch all_comp_08022026. The last quota day is partial.

Figure 6: The quota is the clock. Readings per hour during the February batch, with approximate quota-day totals of 472 to 506 above the bars. The reset boundary is reconstructed, so a day can slightly exceed 500. Source: data/eon.db, frozen in figure-data/evidence.json.

Of the 1,410 companies attempted, 83 failed: 61 because no filing could be retrieved, and 22 after the model call failed repeatedly.

5.4 The database is the batch's memory


A batch that runs for two weeks will be interrupted: by a crash, a reboot, a power cut, a closed laptop. The database is what lets it pick up where it left off without losing or repeating work (Figure 7).

EON uses SQLite, a single file, because the whole system runs on one machine and a server database would add a service to install, configure and keep alive. SQLite's weakness is that only one writer can write at a time. EON works with that rather than against it.

- **Write-ahead logging.** In WAL mode readers never block the writer, so the web interface can browse results while a batch writes new ones.
- **Patience, then retries.** Each connection waits up to 30 seconds for the lock, then retries up to ten times with exponential backoff and random jitter, so that competing workers do not retry in lockstep.
- **Leases.** A worker claims a company with a conditional update (set it to running, but only if it is still pending), so exactly one worker wins. The claim carries a lease of three hours, renewed by a heartbeat every five minutes. If a worker dies, its heartbeat stops, its lease expires, and the company returns to the queue.

The database is the batch's memory

Core tables and the mechanisms that keep many writers from colliding



Also: file_cache, year checkpoints, custom prompts; 14 migrations.

Source: eon/ui/database/repository.py and migrations; eon/ui/services/batch_queue.py.

Figure 7: The database is the batch's memory. Batch tables track who is working on what; analysis tables hold what was read. Beside them, the mechanisms that let many workers and two interfaces share one SQLite file. Sources: eon/ui/database/repository.py and migrations; eon/ui/services/batch_queue.py.

- **Crash recovery.** On start-up EON checks whether the process that last held the batch is still alive. If not, its companies go back to pending.
- **Per-year checkpoints.** Each fiscal year's reading is saved the moment it arrives. An interrupted company resumes from its last completed year, not from the beginning.
- **A unique index.** In February a double-save pattern (once as each year finished, once at the end) quietly stored some readings twice; `INSERT OR IGNORE` had nothing to ignore against. Migration 14 removed the duplicates and added a unique index. (Separate batches that overlapped still read some company-years twice, which is why 6,973 stored rows hold 6,653 distinct company-years; the evaluation keeps the first.)
- **Backups.** During a batch, EON copies the database with SQLite's own backup API and keeps the last five.

The schema itself has been through 14 migrations, each written to be safe to run twice and applied automatically at start-up. One earlier file in the repository's history is named `fcntl.db.corrupted`: a reminder of why the rest exists.

5.5 Two front doors and a doorbell

EON has a command line and a web interface, and they do different jobs.

The command line is for long work. A batch of a thousand companies over ten years needs twenty days of quota; the README's advice is to start it in a terminal multiplexer and detach, because a batch must outlive any browser tab. `eon batch`, `eon analyze`, `eon export` and `eon scan-contrarian` are built with Click.

The Streamlit web interface is for looking: launching a single analysis, browsing history, reading results, managing the batch queue, and resuming anything interrupted. For a while the two had their own code paths and drifted apart. On 6 February both were rebuilt on one shared service layer, so a batch started from either behaves the same way, and a batch started from the command line can be watched from the browser.

The doorbell is Discord. A batch that runs overnight for two weeks needs a way to say something while its owner is asleep or away. EON posts to a Discord webhook when a batch completes or fails, when every key is spent and it is waiting for the reset, and when disk space runs low: the 27.5 GB of PDFs are a real risk. Before starting, it checks that there is room.

5.6 Two machines

EON ran on two computers. The February batch ran on Windows; the database still stores its paths with backslashes. Development happened on a Mac, which kept its own database of 1,184 readings. A pull request titled "multi-os support", merged on 2 February 2026, made the difference invisible. File locks use `portalocker`, which speaks both Windows and Unix locking. Paths are built with the platform's own separators. Stray Chrome processes are cleaned up with `pkill` on the Mac and `taskkill` on Windows.

5.7 Is it efficient?

It depends on what is scarce.

Against the quota, EON is efficient: every request is a whole filing, nothing is sent twice, and nothing is lost to a crash. Against the clock, it is not: the batch spent about 80% of its time waiting. Against compute, it is wasteful but cheap: a browser per filing and about 6 seconds of extraction buy a faithful text. Against storage, it is heavy: 27.5 GB of PDFs for a database of 240 MB.

If requests were plentiful, the design would change. Parsing HTML directly would drop the browser. Caching extracted text instead of PDFs would cut the disk. Sending only the sections each lens needs would cut tokens substantially. Every one of those trades faithfulness for speed, and while the quota was the constraint, there was no reason to make them.

6. Workflows as hypotheses

With the panel built, the questions became narrower. From May 2026 new analyses were written as custom workflows: Python files, discovered automatically, each with its own prompt and schema. They became a way to write down a hypothesis and test it on hundreds of filings.

CSPP turned a long analytical framework (17 scored components in 5 domains) into a reading. Its first version asked for everything in one call, with a schema of about 106 KB, 28 times the size of one lens. When any of its roughly 200 required fields drifted, validation failed and the whole reading was lost. On 20 May 2026 it was split into four smaller calls, merged in code, with a partial result kept if one part fails. The rewrite also stopped trusting the model's arithmetic: it supplies component scores, and Python adds them up.

Moonshot × **Excellence** asked whether a company offers an asymmetric bet and whether it resembles the great companies of 2025. It uses one call on purpose: it can include up to 250,000 tokens of reference material, and a second call would resend all of it and spend another request. It rates six dimensions on anchored 0–10 scales and computes the totals in code. It also carries a watchdog, added after a large filing hung a call for about two hours without an error.

Contrarian Hunter and the **options scans** brought the 2025 experiments into EON. The February 2026 options scan, Asymmetric Options V4, read the latest filing of 1,415 companies in two days and assigned each a directional bias. Section 8.5 tests it.

Two rules came out of this work. Split a question the model cannot answer reliably in one piece. And let the model judge while code does the counting.


Part II · What it read

7. Every question it was asked

Across five versions, the model gave 29,375 stored answers (Figure 8).

29,375 answers in seventeen months

Every stored model answer, by kind of question / bubble area is the number of answers



Sources: origin project files; data/archive/fintel.db; data/eon_mac.db; data/eon.db.

Figure 8: 29,375 answers in seventeen months. Every stored model answer, grouped by the question it answered. Bubble area is the number of answers. Sources: the origin project's files; data/archive/fintel.db; data/eon_mac.db; data/eon.db.

Analysis	When	Answers	Asked
Excellence study	May 2025	15,632	What made great companies great, and who resembles them?
Contrarian scan	May 2025	1,933	Where is strength hidden that others miss?
Options ideas	June 2025	1,933	Buy calls, buy puts, or neither?
Fintel analyses	Dec 2025–Jan 2026	286	Fundamentals, single lenses, syntheses
Three-lens verdicts	Jan–Jul 2026	8,157	Buy, hold or sell, as Buffett, Taleb and a contrarian
Asymmetric Options V4	Feb 2026	1,415	Which way is the tail risk?
CSPP and Moonshot	May–Jun 2026	19	Frameworks too large for one call

Read together, the answers show a model with habits.

- **It likes to hold.** Of 6,411 dated three-lens verdicts, 50.8% say HOLD, 27.5% BUY, 16.5% SELL, 3.8% STRONG BUY and 1.4% STRONG SELL. The mix barely moves from year to year, through a boom, a bear market and a recovery. Conviction is Medium 4,632 times, High 1,516 times and Low only 148 times.
- **It finds danger everywhere when asked to look for it.** Asked where the tail risk lies, the February 2026 options scan gave 873 companies (62%) a put bias and 27 (2%) a call bias.
- **It favours the famous.** The 2025 contrarian scan, meant to find hidden gems, put MicroStrategy first, followed by Apple, AST SpaceMobile, Tesla, Microsoft and Nvidia. Only 27 of 1,933 companies scored 70 or more.
- **Its numbers cluster.** A quarter of all 2025 compounder scores are exactly 68, and three values cover 44% of companies.

None of this says whether the answers were right. That takes a test, and the test is harder than it looks.

8. Was it right?

8.1 Three dates

In February 2026, while the market-wide batch was still running, I checked the readings against stock returns. BUY-rated filings had beaten SELL-rated ones by 14.1 percentage points over two years. The report concluded that the model had genuine stock-picking skill.

It had skipped a question. The model was being graded on years it had already read about. Gemini 2.5 Flash's training data ends in January 2025, according to Google, and this paper uses 31 January 2025 as the line. Show it a 2021 annual report and ask whether to buy: the report ends in 2021, but the model knows how 2022 went. A backtest like that is a trader handed tomorrow's newspaper. Their record would be extraordinary, and it would say nothing about their judgement.

Every verdict therefore has three dates: when the company filed, where the model's knowledge ends, and when the verdict was written. Their order decides what a test can prove (Figure 9).

Every test below uses one reading per company and year; entry at the first close after the filing (or, for the ledgers, after the reading); returns in excess of SPY, the S&P 500 fund; and significance from shuffling verdicts 10,000 times within each year, and within each year and industry. Appendix A has the details.

8.2 With hindsight, and after the cutoff

For one-year windows that closed before the cutoff, 1,099 BUY readings beat SPY by 2.2 points and 657 SELL readings trailed it by 10.9: a spread of 13.1 percentage points ($p = 0.0001$). About 42% of it came from favouring the right industries.

For the 1,443 filings published after the cutoff, which the model cannot have read, the same test gives a different picture (Figure 10). Ranked by where each stock finished within its year, the spread halves, from 11.1 to 6.2 percentile points. The average spread, 14.9 points, looks unchanged, but its uncertainty is ± 23 points.

Where the edge lived also changed (Figure 11). Before the cutoff, the five verdicts formed a perfect staircase: STRONG SELL companies finished at the 43rd percentile of their year, SELL at the 43rd, HOLD at the 50th, BUY at the 54th, STRONG BUY at the 60th. After it, BUY still finished above the middle, but SELL rose to the 49th percentile, level with HOLD. The SELL calls, which had carried most of the historical spread, stopped working.

Three dates decide what a test can prove

When the filing appeared, where the model's knowledge ends, and when the verdict was written

WITH HINDSIGHT



The model may already know how the story ended.

AFTER THE CUTOFF



The model cannot have read the filing, but the verdict came late.

WRITTEN IN ADVANCE



The verdict exists before anything it predicts. Only this is a forecast.

Event order is schematic; spacing is not elapsed time. Source: the evaluation protocols.

Figure 9: Three dates decide what a test can prove. Only the third arrangement, a verdict recorded before anything it predicts, is a forecast in the ordinary sense. Event order is schematic. Source: the evaluation protocols.

A model that remembered which companies collapsed would look like this. So would a model whose caution suited 2021 to 2024 and not 2025. The data cannot choose between them.

What survives the cutoff is real but small. By rank, BUY readings beat SELL readings by **5.9 percentile points** ($p = 0.0059$) within industries. That statistic was chosen after the next figure was seen, and it rests largely on one filing season. By average, the result is inconclusive: 4.6 points at six months ($p = 0.094$) and 14.9 at a year ($p = 0.215$).


8.3 One stock, and twenty-four

The average fails because of one company (Figure 12). Babcock & Wilcox filed its 10-K on 31 March 2025, closed at \$0.46 the next day, and stood at \$15.72 a year later: 3,314 points ahead of SPY. That single reading adds 10 points to the BUY group's average. The price series was checked against Nasdaq's independent record; there was no split. EON read the filing on 9 February 2026, rated it BUY, and by then the stock had already risen about twentyfold. The model could not see that: its training ended first, and its call path has web search switched off.

Averages also hide what the verdicts look like one company at a time. Figure 13 shows 24 companies drawn at random, with a published seed, from the 195 that EON read in every year from 2020 to 2024 and rated both BUY and SELL at least once. Of their 70 BUY and SELL calls, 40 went the right way: 57%. Read row by row, the grid looks less like foresight than a slightly loaded coin.

After the cutoff, ranks halve; averages turn to noise

One-year BUY minus SELL, the same test on both sides of the cutoff / 95% intervals



Before: 1,099 BUY, 657 SELL. After: 339 BUY, 201 SELL.

Source: evaluation/results.json. Intervals: 1.96 standard errors assuming independent readings.

Figure 10: After the cutoff, ranks halve; averages turn to noise. The same one-year BUY-minus-SELL comparison on both sides of the model's knowledge cutoff, with 95% intervals from independent-reading standard errors. The permutation tests in the text are the formal results. Source: evaluation/results.json.

8.4 A ledger written in 2025

The strictest evidence is the oldest. The 2025 scripts wrote their scores to disk in May and June 2025, after the model's cutoff and before any of the returns they might predict. They are a ledger, and they need no argument about training data (Figure 14). Their dates come from file timestamps, not a sealed archive, and I had looked at the first result before writing the test.

Resemblance to great companies pointed the wrong way. Across 1,766 companies, the higher the 2025 compounder score, the worse the next year: Spearman -0.14 , $p = 0.0001$. The top fifth trailed SPY by 17.8 points on average; the bottom fifth beat it by 18.6. A sensible idea, carefully applied, failed its first year. One plausible reason, untested, is that looking like a past winner was already in the price. The original ambition was about decades, which one year cannot settle.

The contrarian score and the options calls showed nothing. Spearman $+0.01$ ($p = 0.71$) for the contrarian score; calls minus puts $+2.8$ percentile points ($p = 0.54$) over six months for the options ideas.


These are not EON's verdicts. Across 1,054 companies scored by both, the compounder score and EON's verdict are almost unrelated (Spearman 0.05).

8.5 A ledger written in February 2026

EON left a ledger of its own. The Asymmetric Options V4 scan of 25–26 February 2026 recorded, for 1,415 companies, a directional bias and an asymmetry score, six months before the outcomes it would be graded on. I wrote the test before running it (Figure 15).

After the cutoff, the SELL calls stop working

Average one-year return percentile by verdict / 0.5 is the middle of the pack



Source: evaluation/stories.json. Numbers: readings. Bars: ± 1.96 SE. Hollow: fewer than 60.

Figure 11: After the cutoff, the SELL calls stop working. Average within-year percentile of one-year excess return by verdict, ± 1.96 standard errors; hollow points have fewer than 60 readings. The verdict order still correlates with outcomes after the cutoff (Spearman 0.094, $p = 0.0012$), driven by BUY rather than SELL. Descriptive. Source: evaluation/stories.json.

On direction, the scan had no skill. Stocks marked for calls did no better than stocks marked for puts (-2.9 percentile points, $p = 0.62$); put-bias stocks in fact beat SPY slightly more often than the rest. On magnitude, it did. Its asymmetry score correlated with the size of the move, up or down (Spearman 0.12, $p = 0.0001$), and stocks marked “straddle”, the model’s way of saying something big will happen, moved more than those marked “no edge” ($p = 0.031$).


This is the cleanest positive result in the project, and it is the modest kind. A model reading a 10-K can tell which companies are fragile, levered, dependent on a binary event, or opaque, and those companies move more. It cannot tell which way. That is what an options trader would call volatility, and it is already in the price of options.

8.6 What the tests say together

Graded with hindsight, the model looked like an analyst. Graded on what it could not have known, it looked like a careful reader with a small, fading edge, one bad idea and one modest talent. The evidence has limits that apply throughout: every clean result comes from roughly the same eighteen months; the universe was drawn from companies still listed in 2026; industry is controlled but size, value, quality and momentum are not; extraction and the model’s facts are unaudited; and returns are not trades. Section 11 is what the paper does about that.

One stock can decide a mean

Post-cutoff filings / one-year excess over SPY, percentage points / symmetric log axis



Solid line: mean. Dotted line: median. Linear between -100 and $+100$ points.

Source: evaluation/trades.csv; BW checked against Nasdaq closes. Filed after 31 Jan 2025.

Figure 12: One stock can decide a mean. Every post-cutoff BUY and SELL reading with a one-year outcome. Solid lines are means, dotted lines medians; the axis is linear to ± 100 points and logarithmic beyond. Sources: evaluation/trades.csv; evaluation/verification/bw-price-check.json.

Company by company, the edge is thin

24 companies drawn at random, seed 20260930 / one-year excess return over SPY, points

							cutoff
		FY2020	FY2021	FY2022	FY2023	FY2024	
ACIW	ACI Worldwide	-27 ✓	-17	-8	+54	-44 X	
BW	Babcock & Wilcox Inc	-9	-19 ✓	-118	+4	+3,314 ✓	
COKE	Coca-Cola Consolidated	+80 X	+32 ✓	+24 ✓	+53 ✓	+26 ✓	
CTBI	Community Trust Bancorp	-9 X	+12	-35	+23 ✓	-2 ✓	
DGX	Quest Diagnostics	+1 ✓	+14 ✓	-39 X	+21 ✓	+8 X	
EME	EMCOR Group	+4	+52 ✓	+38 X	+9 ✓	+64 ✓	
EVRG	Evergy	+6 ✓	+5 X	-43	+28	+12	
EYPT	EyePoint	-2	-75	+724	-79 X	+98 X	
FBNC	First Bancorp	-2 X	+2 ✓	-48	+6 X	+25 X	
FMBH	First Mid Bancshares	-12 X	-17	-25 ✓	+8	-7 X	
GNTX	Gentex	-27 X	+2 ✓	+1	-50	-19 ✓	
IEP	Icahn	-11 X	+22 X	-86 ✓	-54	-19	
KTOS	Kratos Defense & Secur...	-49	-9	+14 X	+6 ✓	+227	
MGRC	McGrath RentCorp	-16 X	+38 X	-9 X	-17	-24	
MPB	Mid Penn Bancorp	-8 X	+9 X	-56 X	+25 ✓	+5 ✓	
MTZ	MasTec	-25	+38 ✓	-35 ✓	+22 X	+119	
MYRG	MYR Group	+41 ✓	+36	+18	-41	+109 X	
NIC	Nicolet Bankshares	+11 ✓	-13 ✓	-23	+34	+15	
NWBI	Northwest Bancshares	-9 X	+12	-41 ✓	-1	-10	
OC	Owens Corning	+8 ✓	+9 ✓	+18 ✓	+1	-39 ✓	
PRLB	Proto Labs	-80 ✓	-37	-16	-10	+38 ✓	
UVE	Universal Insurance Ho...	-31 ✓	+82 X	-20 X	-4	+50	
VLV	Valley National Bancorp	+1 ✓	-5 ✓	-60	+6	+18	
ZD	Ziff Davis	-13	-17	-43	-60 ✓	-54 X	

Strong buy Buy Hold Sell Strong sell ✓ 40 of 70 calls

Source: evaluation/stories.json. Pool: 195 companies read every year, with a BUY and a SELL.

Figure 13: Company by company, the edge is thin. Colour is the verdict; the number is the one-year return over SPY in percentage points; a tick marks a BUY that beat SPY or a SELL that trailed it. The rule selects on verdicts, never outcomes. Because the median stock trailed SPY in these years, SELL ticks come cheaply. Source: evaluation/stories.json.

Looking like a great company was a bad sign

Scores written down in 2025, before their outcomes / top minus bottom, return percentile




Band: middle 95% of industry shuffles, not a confidence interval.

Source: evaluation/origin-ledger/results.json. Top/bottom score groups include threshold ties.

Figure 14: Looking like a great company was a bad sign. Top fifth minus bottom fifth of each 2025 score (calls minus puts for options), in percentile points of return, against the middle 95% of 10,000 industry-preserving shuffles. Source: evaluation/origin-ledger/results.json.

It saw how far a stock would move, not which way

Options scan of 25–26 February 2026, graded six months on / written before the outcome



Source: evaluation/stories.json; 1,365 companies, 26 with a call bias.

Figure 15: It saw how far a stock would move, not which way. Six-month outcomes for 1,365 companies, by the bias the scan recorded in February 2026. Left: direction, as the average percentile of excess return. Right: magnitude, as the average percentile of the absolute excess return. Source: evaluation/stories.json.


Part III · What it taught

9. How the model went wrong

The most useful findings in this project are not about stocks. They are about how a fluent language model fails when it is asked to do a job thousands of times (Figure 16).

Eight ways a fluent model went wrong

Measured in the project's own outputs, 2025–2026



Sources: origin files; evaluation outputs; data/eon.db; the flattening script's notes (~130).

Figure 16: Eight ways a fluent model went wrong. Each number is measured in the project's own stored outputs. Sources: the origin files; evaluation outputs; data/eon.db; the flattening script's notes.

It did not know what day it was. Run in May 2025, the model dated 1,916 of 1,919 analyses 2024, a year behind, and proposed 91 option expiries that had already passed. It reasons from its training-time present unless told otherwise. EON's later workflows stamp dates in code.

It may remember the future. A model trained after the events it is asked to judge can borrow from them without saying so. A search of every reading for terms with known first-use dates found no careless leaks: the Inflation Reduction Act appears in 300 readings, never in one of a filing written before the law. But a verdict nudged by memory leaves no trace in its prose.

Its numbers are less precise than they look. A quarter of all compounder scores were exactly 68. A 0–100 scale that behaves like a four-point one invites false precision. The later workflows use anchored scales, each defined at 0, 5 and 10.

It leans. It holds half the time. Asked to look for tail risk, it found it in 62% of companies and upside in 2%. The framing of a question shapes the answer's direction as much as the evidence does.

It favours platitudes and the famous. It credited “continuous innovation” to every great company, and its contrarian scan's hidden gems included Apple, Microsoft and Nvidia.

It drifts in format. Without a schema, six factors took 130 spellings. With one, 462 verdicts still opened with a label (“Overall investment recommendation: ...”) instead of the answer, and conviction arrived in half a dozen phrasings. Every parser must expect this.

It cannot always add. CSPP stopped trusting its arithmetic and recomputed every total in code.

It sometimes stops. One large filing hung a call for about two hours without an error. Every call now has a watchdog.

It values without prices. Apple's SELL rests on a “highly stretched valuation”, but a 10-K contains no current share price. The model supplies one from memory, which is exactly the channel through which hindsight can leak.

10. Why it was worth doing

EON was never going to rival Citadel, Point72 or the Medallion Fund. Those firms buy data nobody else has, trade in milliseconds, employ hundreds of researchers, and move capital at scale. Medallion is famous for short-horizon statistical patterns, not opinions about annual reports. A 10-K is the most public document a company produces. Thousands of professionals read the important ones on the day they are filed; whatever it says is in the price within hours. A machine reading the same text a year later, on free quotas, was not going to find money lying on the pavement.

It was worth doing anyway, because it was three projects at once.

As a finance project, it turned a vague question (can reading annual reports well tell you anything?) into a measurable one, and taught the details that make such measurements honest: excess returns and their benchmark, industry composition, survivorship, the difference between an average and a rank, one stock deciding a mean, and above all look-ahead. The answer was humbling in the right way: what a filing says is largely known, and what it signals most reliably is risk, not direction.

As an AI project, it measured a language model at scale, on the same task, thousands of times, and caught it misdating its work, drifting in format, clustering its scores, leaning bearish when asked about risk, and possibly remembering the future. Those are not failures of a particular prompt. They are properties of fluent models, and they only show up in volume.

As a systems project, it forced real engineering under a hard constraint: cross-process locks, key rotation, leases and heartbeats, crash recovery, a database built to survive two weeks of interruption, two operating systems, two interfaces, and an alarm that works while you sleep. None of that was planned. Every piece answers a failure that happened.

And it left something no one else has: 29,375 answers written down with their dates, a method for testing them honestly, and a bet.

11. The bet

Every lesson in this paper points to one practice: write the verdict down, with its date, before the future arrives. So the paper makes a forecast of its own (Figure 17).

The bet: 49 calls to check in October 2027

EON's strongest current verdicts, frozen 28 September 2026 / every one, not a selection

STRONG BUY · 30

AD	AMG	AMPH	BBW	BCRX	CB	CNK	CRK	EME	GRBK
IBKR	INVA	IRDM	KFY	LNG	NEU	NXST	PEGA	PFS	POWL
QCRH	RNR	SBH	SCSC	SRCE	SYF	USLM	VSH	WINA	WLKP

STRONG SELL · 19

AAOI	ALX	ARDX	AVAV	CELH	CSR	CWST	DCOM	INDV	NUTX
OGS	ORC	PLUG	RGEN	RIG	SMR	UHAL	UNIT	ZBRA	

All 1,257 verdicts are frozen in `evaluation/sealed-ledger/ledger.csv`, SHA-256 `b5a03624d885c0c9...`

Primary test fixed in advance: BUY-type minus SELL-type return rank over the following year.

Figure 17: The bet: 49 calls to check in October 2027. EON's latest verdict on each of 1,257 companies, frozen on 28 September 2026 with their prices, and graded one year after publication. The tiles are every STRONG BUY and STRONG SELL, not a selection. Source: `evaluation/sealed-ledger/`.

The ledger holds EON's most recent verdict for each of 1,257 companies: 634 HOLD, 341 BUY, 233 SELL, 30 STRONG BUY and 19 STRONG SELL. The whole file is fixed by its SHA-256 fingerprint, which begins `b5a03624d885c0c9`. The test was chosen in advance: BUY-type minus SELL-type percentile rank of one-year return over SPY, from the first close after 30 September 2026. Many of these verdicts are already a year old, since they read filings from early 2025; that is part of the test.

The next vintage should do better than a frozen file on one laptop. Read the fiscal 2025 reports now, with the same model and prompt; store each reading with the filing's identifier, a hash of the text the model saw, the prompt version and the time; pin the model, because a new model has a new cutoff and turns clean evidence back into hindsight; tell it the date; and publish the fingerprint somewhere public before the returns arrive.

12. Conclusion

I set out to build a patient reader. It took five versions and seventeen months, and most of the work was not reading at all. It was waiting for midnight, printing web pages, keeping two dozen workers from colliding, and making sure nothing was lost when something broke.

The reader works. It is specific, tireless and consistent enough to be counted. Counting it produced a result that looked like foresight, and arranging the evidence by what the model could have known took most of that away. What remains is modest and real: a small edge in its BUY calls, a sense of which companies will move a lot, and a clear view of how a fluent model goes wrong.

The newspaper the model read ends in January 2025. Everything after that has to be read the ordinary way, one day at a time.

Write the question once. Keep the answer with its date. Let the next year answer back.

Appendix A: methods

Scope. One reading per ticker and fiscal year, the first stored. Of 6,653 such readings, 6,651 were made with Gemini 2.5 Flash; two made with Gemini 3.5 Flash are excluded. A usable observation needs a parsed verdict, a filing date and prices: 6,203 readings of 1,309 companies.

Returns. Entry at the first close strictly after the filing date. Excess return is the change in adjusted close over 126, 252 or 504 trading sessions minus SPY's change over the same dates, in percentage points, not annualised. For the ledgers, entry follows the recorded reading date. Prices are a frozen Yahoo Finance snapshot through 28 September 2026; the largest outlier was checked against Nasdaq.

Samples. "Before the cutoff" means the return window closed before 31 January 2025. "After the cutoff" means the filing was published after it.

Spreads and ranks. The mean spread is the BUY-group average excess return minus the SELL-group average. The rank spread uses each reading's percentile within its fiscal year, so no single stock can dominate.

Permutation tests. Verdicts are shuffled 10,000 times within fiscal years, and separately within fiscal-year-and-industry groups (industry from a February 2026 FactSet snapshot; single-reading groups dropped). Two-sided p-values measure distance from the shuffled mean, with a plus-one correction; 0.0001 is the floor.

What was fixed in advance. The historical one-year and post-cutoff six-month tests were specified first. The post-cutoff one-year test, the rank statistic and the anachronism search were added after earlier results were seen. The verdict ladder, company grid and February 2026 options test follow `stories-config.yaml`, written before they were run. None is adjusted for multiple comparisons.

Ledgers. The 2025 scores are read from the origin project without modification; compounder dates come from file modification times, other dates from filenames. The February 2026 options scan uses each reading's stored timestamp.

Pipeline measurements. Page counts, characters and extraction times come from 40 randomly sampled cached 10-K PDFs; markup ratios and XBRL counts from 17 raw EDGAR 10-K files. Tokens are estimated at four characters each.

Appendix B: related work and references

Sarkar and Vafa show that pretrained language models carry information about periods after their nominal analysis date, and test for it with events that should be unpredictable; the anachronism search in Section 9 is a simple instance. Glasserman and Lin measure look-ahead bias in model-scored news sentiment and reduce it by removing company names. Lopez-Lira and Tang find that model readings of headlines predict next-day returns; EON reads a different document over a longer horizon, and its results are not a replication.

- Glasserman, P. and Lin, C. (2023). *Assessing Look-Ahead Bias in Stock Return Predictions Generated by GPT Sentiment Analysis*. arXiv:2309.17322.
- Google. *Gemini 2.5 Flash model documentation*, knowledge cutoff January 2025. ai.google.dev/gemini-api/docs/models/gemini-2.5-flash.
- Lopez-Lira, A. and Tang, Y. (2023). *Can ChatGPT Forecast Stock Price Movements? Return Predictability and Large Language Models*. arXiv:2304.07619.
- Sarkar, S. K. and Vafa, K. (2024). *Lookahead Bias in Pretrained Language Models*. SSRN 4754678.

- US Securities and Exchange Commission. *EDGAR fair access policy*: no more than ten requests per second.

Appendix C: reproduction

The Markdown is the editorial source; LaTeX and PDF are generated from it. `BRIEF.md` sets the editorial standard, `figures.md` documents every figure, `graphics.md` briefs the commissioned illustrations, and `revision-notes.md` records every change and open question.

Path	Role
<code>eon/data/sources/sec/</code>	Download, Chrome conversion, extraction, SEC queue
<code>eon/ai/</code>	Prompts, key manager, request queue, Gemini provider
<code>eon/ui/database/</code>	Repository, migrations v001–v014
<code>eon/ui/services/batch_queue.py</code>	Leases, heartbeats, quota waits, alerts
<code>evaluation-config.yaml</code>	Main evaluation and its amendments
<code>origin-ledger-config.yaml</code>	2025 ledger test
<code>stories-config.yaml</code>	Ladder, company draw, February 2026 options test
<code>sealed-ledger-config.yaml</code>	The bet

From the EON root:

```
python docs/whitepaper/evaluate_backtest.py
python docs/whitepaper/evaluate_origin_ledger.py
python docs/whitepaper/evaluate_stories.py
.venv/bin/python docs/whitepaper/measure_pipeline.py
python docs/whitepaper/prepare_figure_data.py
python docs/whitepaper/render_figures.py
python docs/whitepaper/validate_paper.py
python docs/whitepaper/build_paper.py --pandoc /path/to/pandoc
cd docs/whitepaper && tectonic -X compile eon-whitepaper.tex
```

The evaluators open databases read-only and use saved prices; the origin evaluator reads `stock_stuff_06042025/10K_automator` without changing it. The sealed ledger's script refuses to run twice.

Prepared from the EON repository and its predecessors, their databases and files, and a frozen price snapshot. Negative and inconclusive results are retained. Research, not investment advice.